

Structures, Alignment, Unions

Topics

- **Structs in C and assembly layout**
- **Alignment**
- **Unions**

Announcements

- **Bomb Lab #2 update – extension to Friday March 3 by 11:55 pm**
 - We are planning to release an alternate bomb server
 - Email any issues to the TAs and Prof. Han
- **Read Chapter 3.1-3.12 (except 3.11) and do practice problems**

Announcements

- **Midterm #1 Tuesday Feb 28**
 - In class
 - Covers Chapters 1, 2.1-2.3 (no floating point), 3.1-3.9 (no buffer overflow or floating point)
 - Write your answers on the printed exam, then upload them in last 15 minutes to the moodle – bring your laptop!
 - Also turn in paper copy as backup
 - No electronics except during the last 15 minutes of exam
 - You may not use any other online resource except the moodle during the upload time.
 - The moodle will cut off exam submissions at the end of class, so make sure your answers are entered before the end of class!
 - A password to access the midterm on the moodle will be provided in class

Announcements

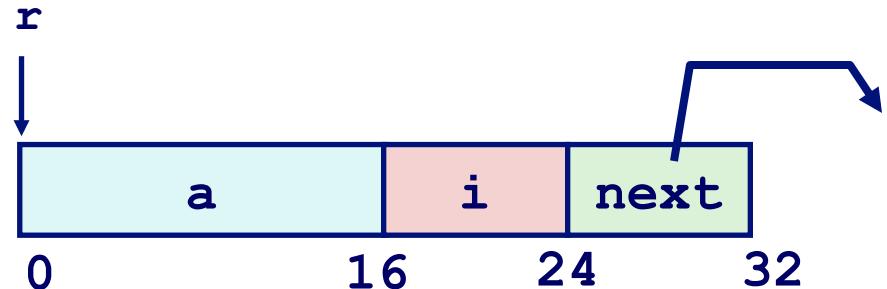
- **Midterm #1 Tuesday Feb 28**
 - Closed book but can bring 1 page summary sheet front & back – write anything you want on it
 - We'll provide a packet of tables from Chapter 3 – print these out and bring them with you
 - TAs may hold extra office hours next Monday
 - We will set up a separate room for those who need extended time for the midterm – contact your TA or the professor to determine when and where to show up (different location than the main lecture hall)
 - Make sure your laptop is charged and functioning properly

Announcements

- **Midterm #1 Tuesday Feb 28**
 - We'll release a practice exam from a previous year later this week - note this was on 32-bit stack conventions, not 64-bit
 - Kinds of questions:
 - Bit manipulation
 - Addition and overflow
 - Call stack
 - C-assembly fill-in-the-blank
 - Good study material:
 - Textbook and its practice problems
 - Lecture slides
 - Practice exam
 - Quizzes

Structure Representation

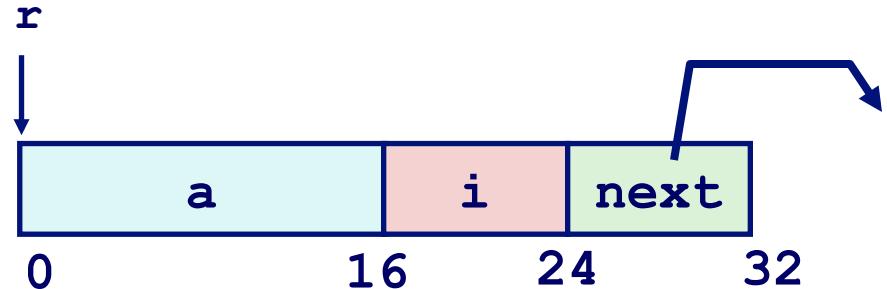
```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
} r;
```



- **Structure represented as block of memory**
 - Big enough to hold all of the fields
- **Fields ordered according to declaration**
 - Even if another ordering could yield a more compact representation
- **Compiler determines overall size + positions of fields**
 - Machine-level program has no understanding of the structures in the source code

Structure Representation (2)

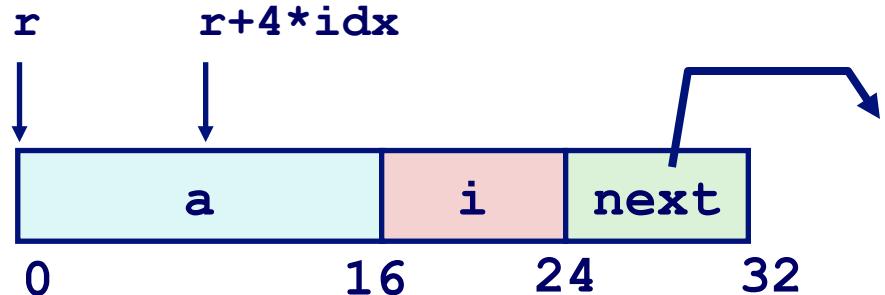
```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
} r;
```



- Refer to members within structure by names
- Members may be of different types
- Accessing/setting a value of a structure member
 - `r.i = 7;`
 - `r.a[0] = 21;`
- Suppose `q` is declared as a “`struct rec *q`”, then
 - `r.next = q;`
 - `q->i = 4; // equivalent to (*q).i = 4`

Generating Pointer to Structure Member

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
} r;
```



- Generating Pointer to Array Element
 - Offset of each structure member determined at compile time
 - Compute as `r + 4*idx`

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```

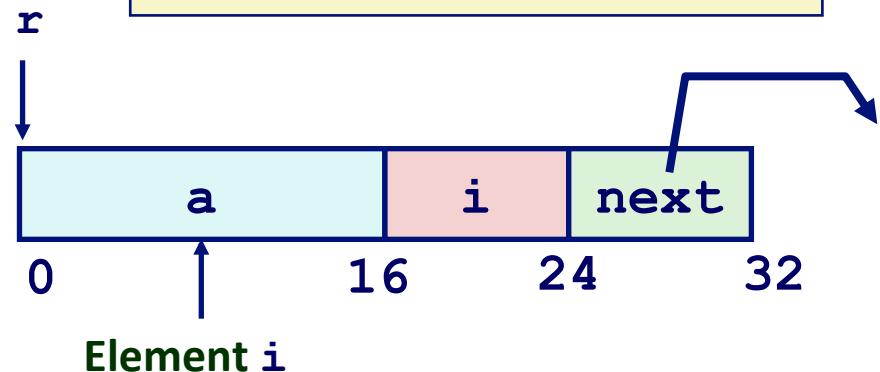
```
# r in %rdi, idx in %rsi  
leaq (%rdi,%rsi,4), %rax  
ret
```

Following Linked List

- C Code

```
void set_val
    (struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movslq  16(%rdi), %rax      #   i = M[r+16]
    movl    %esi, (%rdi,%rax,4) #   M[r+4*i] = val
    movq    24(%rdi), %rdi      #   r = M[r+24]
    testq   %rdi, %rdi         #   Test r
    jne     .L11                #   if !=0 goto loop
```

Alignment

Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K
- Required on some machines; advised on x86-64
 - treated differently by Linux and Windows!

Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (double or quad-words) – system-dependent
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory very tricky when datum spans 2 pages

Compiler

- Inserts gaps in structure to ensure correct alignment of fields

Specific Cases of Alignment (x86-64)

Size of Primitive Data Type:

- 1 byte (e.g., `char`)
 - no restrictions on address
- 2 bytes (e.g., `short`)
 - lowest 1 bit of address must be 0_2
- 4 bytes (e.g., `int`, `float`)
 - lowest 2 bits of address must be 00_2
- 8 bytes (e.g., `double`, `long`, `char *`, ...)
 - lowest 3 bits of address must be 000_2
- 16 bytes (`long double`)
 - GCC on Linux:
 - » Lowest 4 bits of address must be 0000_2

Satisfying Alignment with Structures

Offsets Within Structure

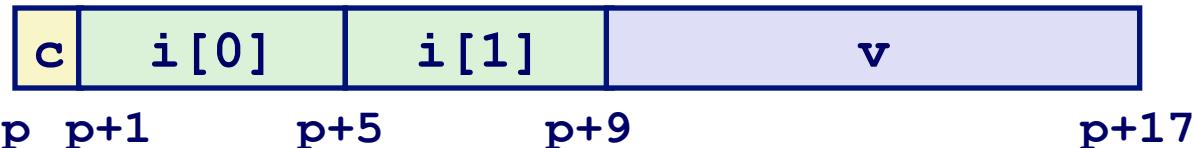
- Must satisfy element's alignment requirement

Overall Structure Placement

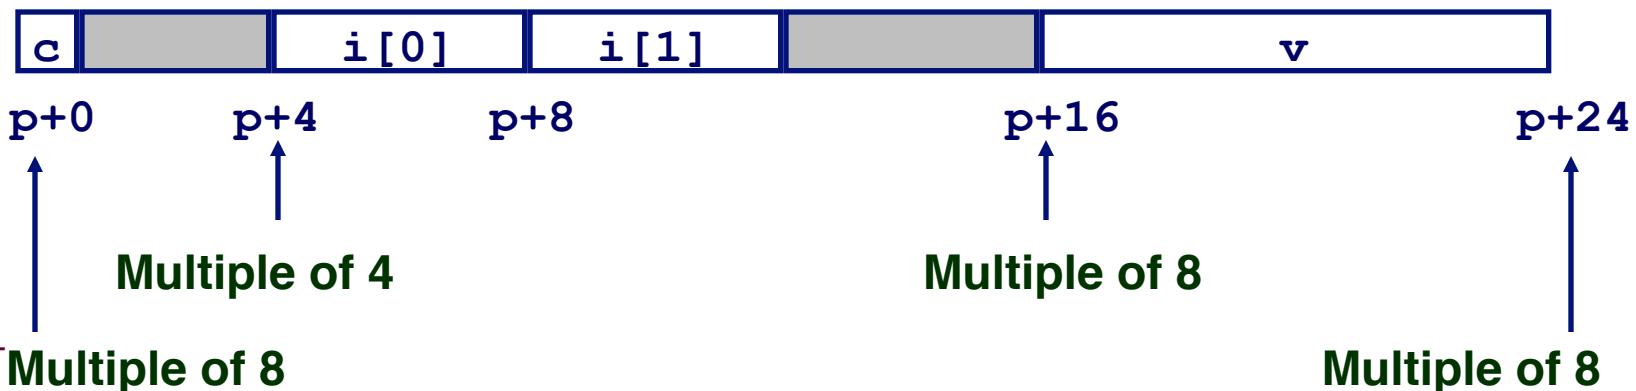
- Each structure has alignment requirement K
 - Largest alignment of any element
- Initial address must be multiple of K
- structure length must be multiple of K

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

Unaligned data:



Aligned data: (K=8 due to double)

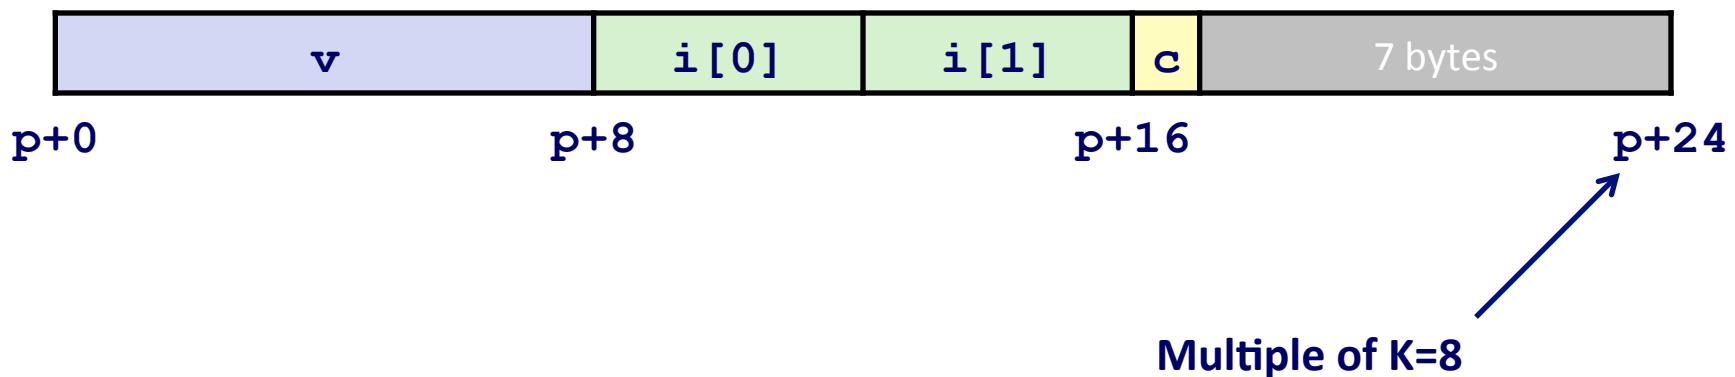


Meeting Overall Alignment Requirement

For largest alignment requirement K

Overall structure must be multiple of K

```
struct s2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```

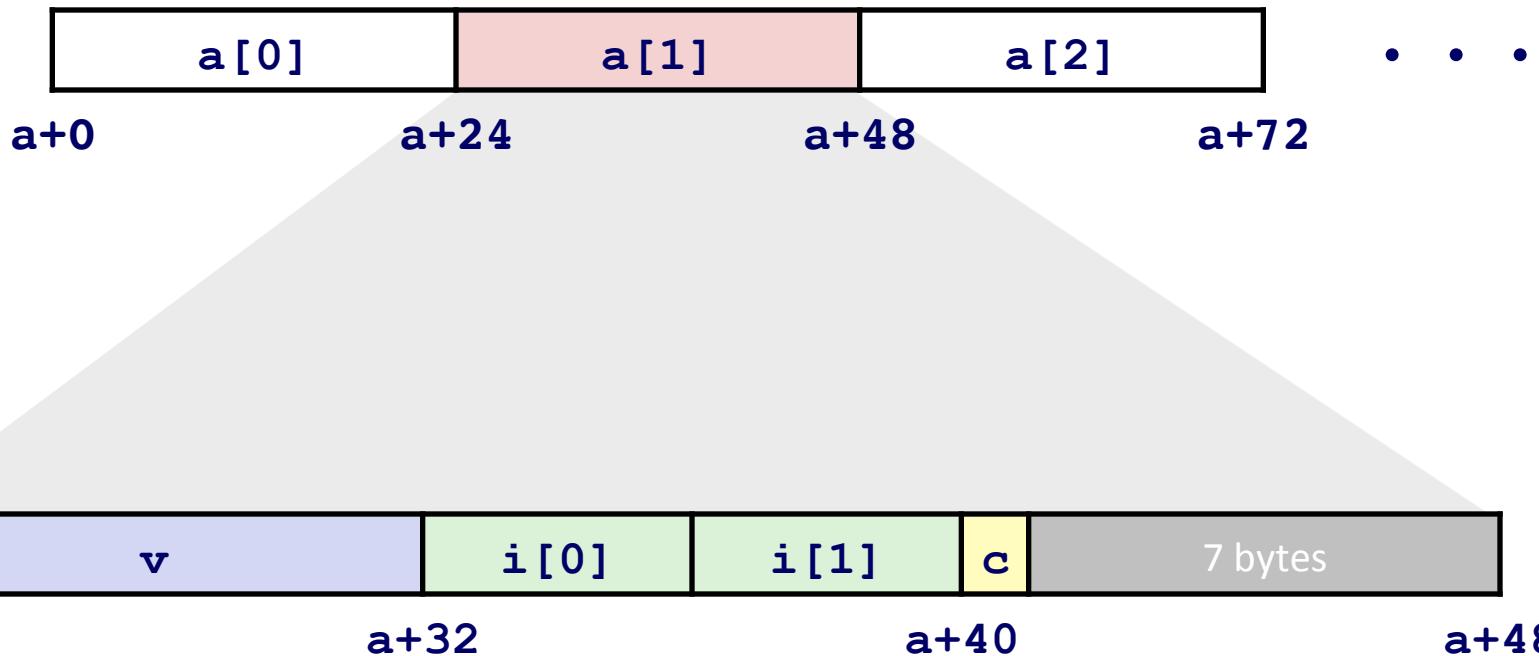


Arrays of Structures

Overall structure length multiple of K

Satisfy alignment requirement
for every element

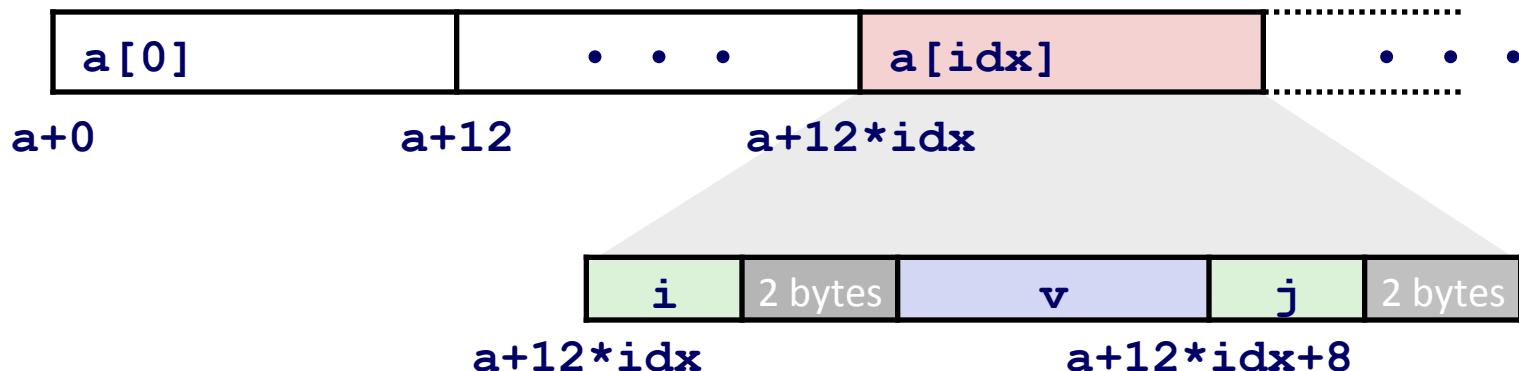
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Array Elements

Compute array offset $12 * \text{idx}$

- `sizeof(S3)`, including alignment spacers



```
short get_j(int idx)
{
    return a[idx].j;
}
```

```
# %rdi = idx
leaq (%rdi,%rdi,2),%rax # 3*idx
movzwl a+8(%rax,4),%eax
```

Element `j` is at offset 8 within structure

Assembler gives offset `a+8`

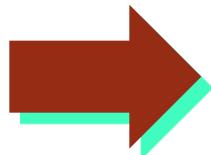
- Resolved during linking

```
struct S3 {
    short i;
    float v;
    short j;
} a[10];
```

Saving Space

Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

Effect (K=4)

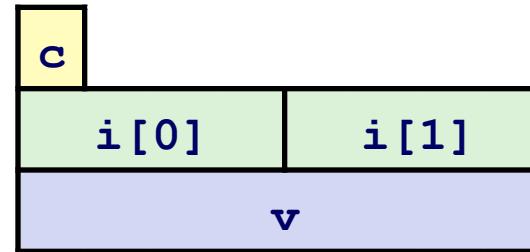


Union Allocation

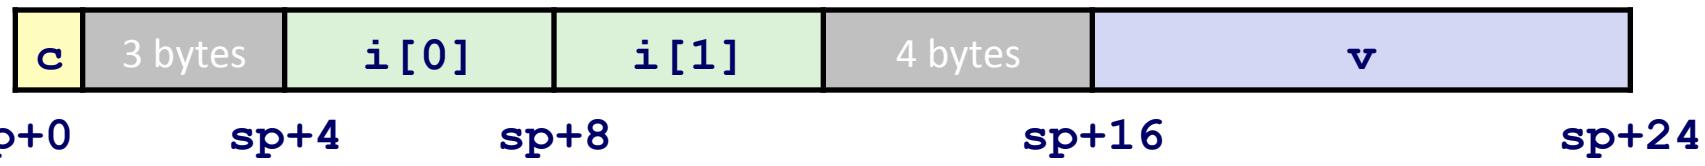
- Allocate according to largest element – overlay union elements
- Can only use one field at a time

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```

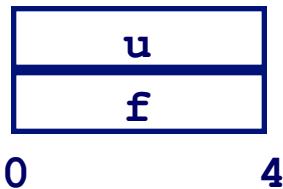


up+0 up+4 up+8



Using Unions to Access Bit Patterns

```
typedef union {
    float f;
    unsigned u;
} bit_float_t;
```



```
float bit2float(unsigned u)
{
    bit_float_t arg;
    arg.u = u;
    return arg.f;
}
```

```
unsigned float2bit(float f)
{
    bit_float_t arg;
    arg.f = f;
    return arg.u;
}
```

- Get direct access to bit representation of float
- bit2float generates float with given bit pattern
 - NOT the same as (float) u converts unsigned int to IEEE fp and rounds, see page 120 of text

- float2bit generates bit pattern from float

Byte Ordering Revisited

Idea

- Short/long/quad words stored in memory as 2/4/8 consecutive bytes
- Which byte is most (least) significant?
- Can cause problems when exchanging binary data between machines

Big Endian

- Most significant byte has lowest address
- Sparc

Little Endian

- Least significant byte has lowest address
- Intel x86, ARM Android and IOS

Bi Endian

- Can be configured either way
- ARM

Using Unions to Study Byte Ordering

```
union {
    unsigned char c[8];
    unsigned short s[4];
    unsigned int i[2];
    unsigned long l[1];
} dw;
```

32-bit	c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
	s[0]		s[1]		s[2]		s[3]	
	i[0]				i[1]			
	l[0]							

64-bit	c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
	s[0]		s[1]		s[2]		s[3]	
	i[0]				i[1]			
	l[0]							

Byte Ordering Example (Cont).

```
int j;
for (j = 0; j < 8; j++)
dw.c[j] = 0xf0 + j;

printf("Characters 0-7 == [0x%x,0x%x,0x%x,0x
%x,0x%x,0x%x,0x%x,0x%x]\n",
dw.c[0], dw.c[1], dw.c[2], dw.c[3],
dw.c[4], dw.c[5], dw.c[6], dw.c[7]);

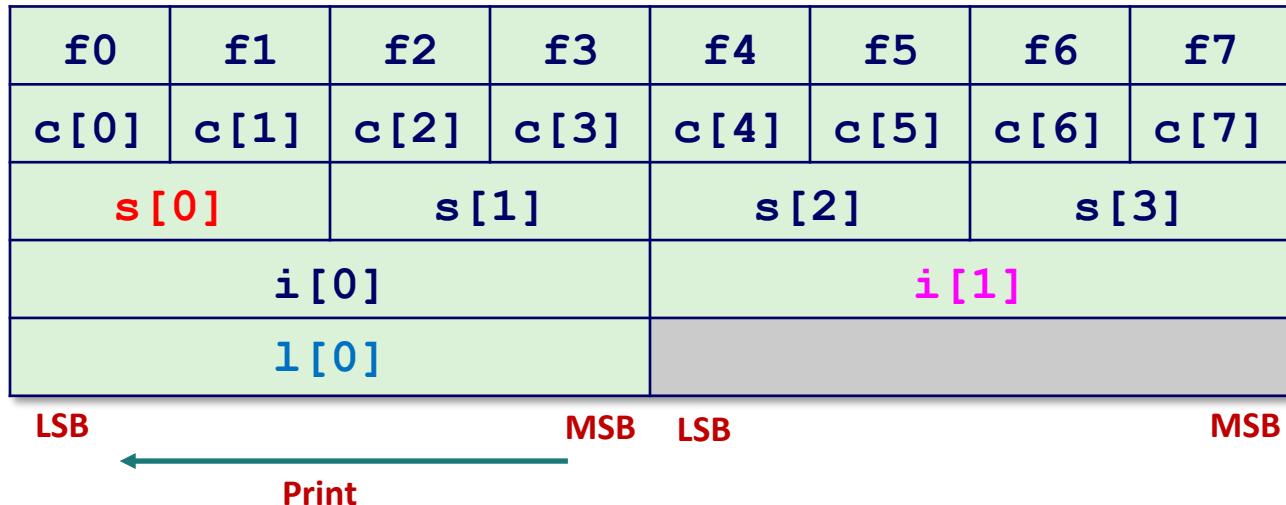
printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x]
\n",
dw.s[0], dw.s[1], dw.s[2], dw.s[3]);

printf("Ints 0-1 == [0x%x,0x%x]\n",
dw.i[0], dw.i[1]);

printf("Long 0 == [0x%lx]\n",
dw.l[0]);
```

Byte Ordering on IA32

Little Endian

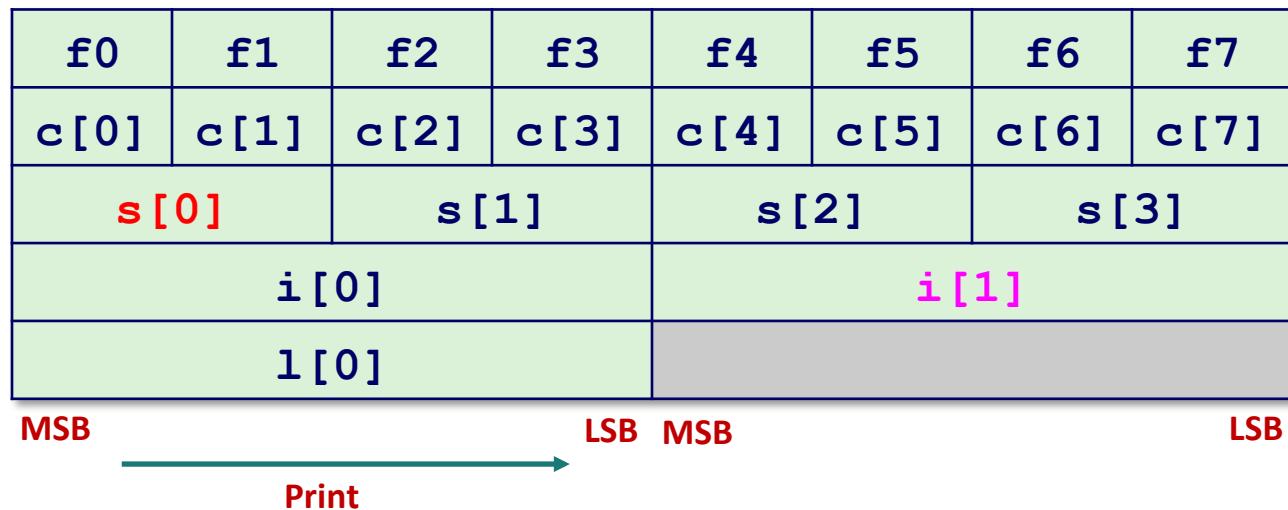


Output:

Characters	0-7	==	[0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts	0-3	==	[0xf1f0,0xf3f2,0xf5f4,0xf7f6]
Ints	0-1	==	[0xf3f2f1f0,0xf7f6f5f4]
Long	0	==	[0xf3f2f1f0]

Byte Ordering on Sun

Big Endian

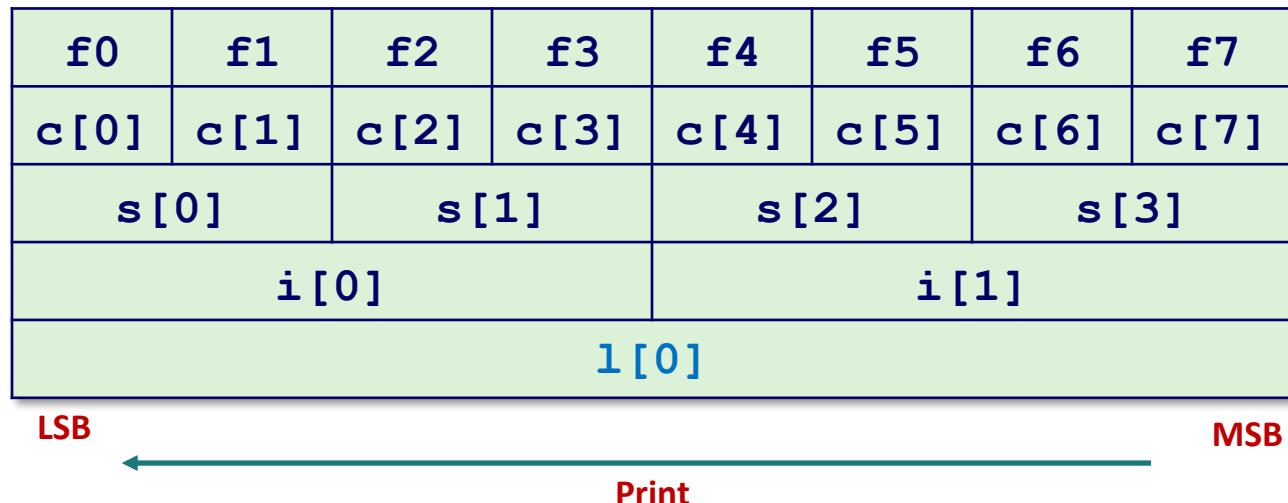


Output on Sun:

Characters	0-7 ==	[0xf0, 0xf1, 0xf2, 0xf3, 0xf4, 0xf5, 0xf6, 0xf7]
Shorts	0-3 ==	[0xf0f1, 0xf2f3, 0xf4f5, 0xf6f7]
Ints	0-1 ==	[0xf0f1f2f3, 0xf4f5f6f7]
Long	0 ==	[0xf0f1f2f3]

Byte Ordering on x86-64

Little Endian



Output on x86-64:

```
Characters 0-7 == [0xf0,0xf1,0xf2,0xf3,0xf4,0xf5,0xf6,0xf7]
Shorts     0-3 == [0xf1f0,0xf3f2,0xf5f4,0xf7f6]
Ints       0-1 == [0xf3f2f1f0,0xf7f6f5f4]
Long        0 == [0xf7f6f5f4f3f2f1f0]
```

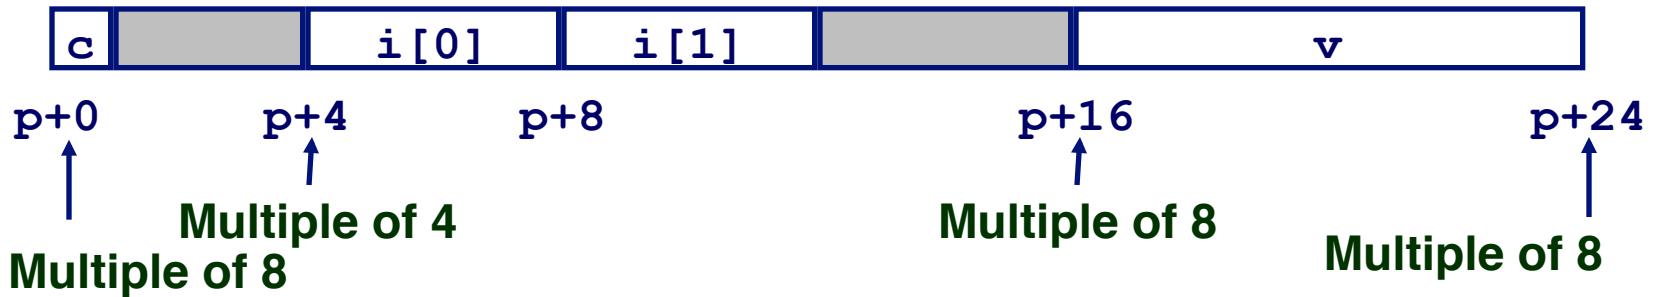
Hence the union is useful not just for determining Endianness,
but also whether this is a 32-bit or 64-bit machine

Supplementary Slides

Linux vs. Windows

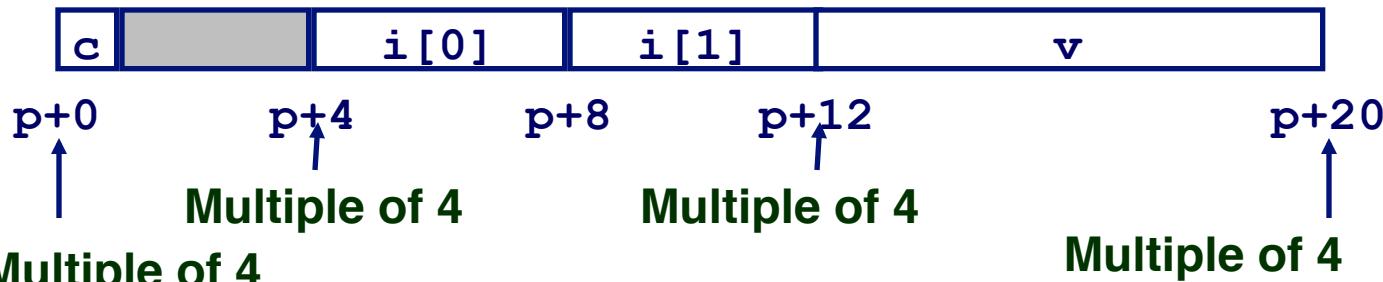
Windows (including Cygwin):

- K = 8, due to double element



Linux:

- K = 4; `double` treated like a 4-byte data type

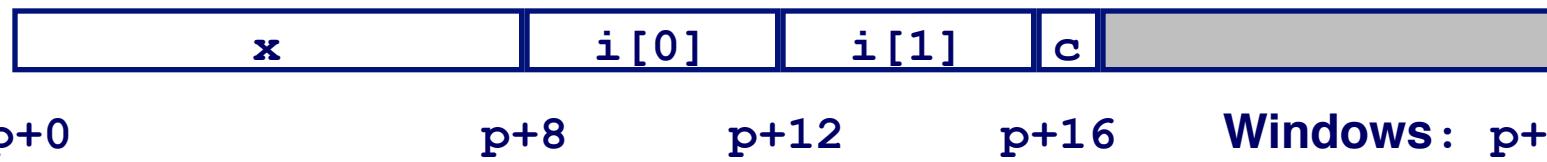


```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

Overall Alignment Requirement

```
struct S2 {  
    double x;  
    int i[2];  
    char c;  
} *p;
```

**p must be multiple of:
8 for Windows
4 for Linux**



```
struct S3 {  
    float x[2];  
    int i[2];  
    char c;  
} *p;
```

p must be multiple of 4 (in either OS)



Ordering Elements Within Structure

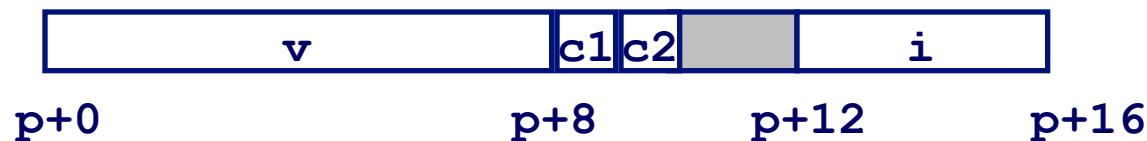
```
struct S4 {  
    char c1;  
    double v;  
    char c2;  
    int i;  
} *p;
```

10 bytes wasted space in Windows



```
struct S5 {  
    double v;  
    char c1;  
    char c2;  
    int i;  
} *p;
```

2 bytes wasted space



Basic Data Types

Integral

- Stored & operated on in general registers
- Signed vs. unsigned depends on instructions used

Intel	GAS	Bytes	C
byte	b	1	[unsigned] char
word	w	2	[unsigned] short
double word	l	4	[unsigned] int

Floating Point

- Stored & operated on in floating point registers

Intel	GAS	Bytes	C
Single	s	4	float
Double	l	8	double
Extended	t	10/12	long double

C pointer declarations

`int *p`

p is a pointer to int

`int *p[13]`

p is an array[13] of pointer to int

`int *(p[13])`

p is an array[13] of pointer to int

`int **p`

p is a pointer to a pointer to an int

`int (*p)[13]`

p is a pointer to an array[13] of int

`int *f()`

f is a function returning a pointer to int

`int (*f)()`

f is a pointer to a function returning int

`int (*(*f())[13])()`

f is a function returning ptr to an array[13] of pointers to functions returning int

`int (*(*x[3])())[5]`

x is an array[3] of pointers to functions returning pointers to array[5] of ints

Understanding Pointers & Arrays #1

Decl	<i>An</i>			<i>*An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>						
<code>int *A2</code>						

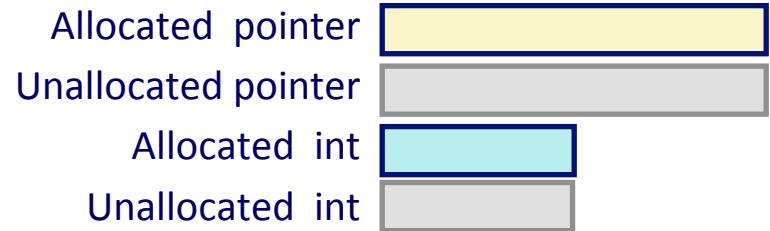
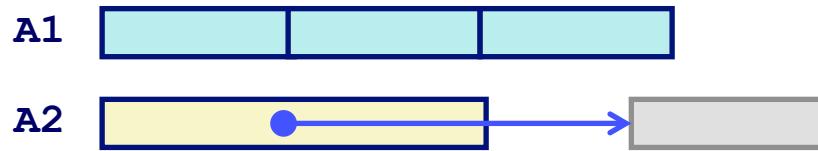
Cmp: Compiles (Y/N)

Bad: Possible bad pointer reference (Y/N)

Size: Value returned by `sizeof`

Understanding Pointers & Arrays #1

Decl	<i>An</i>			<i>*An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size
int A1[3]	Y	N	12	Y	N	4
int *A2	Y	N	8	Y	Y	4



Cmp: Compiles (Y/N)

Bad: Possible bad pointer reference (Y/N)

Size: Value returned by sizeof

Understanding Pointers & Arrays #2

Decl	An			*An			**An		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>									
<code>int *A2[3]</code>									
<code>int (*A3)[3]</code>									
<code>int (*A4[3])</code>									

Cmp: Compiles (Y/N)

Bad: Possible bad pointer reference (Y/N)

Size: Value returned by sizeof

Understanding Pointers & Arrays #2

Decl	An			*An			**An		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4	N	-	-
<code>int *A2[3]</code>	Y	N	24	Y	N	8	Y	Y	4
<code>int (*A3)[3]</code>	Y	N	8	Y	Y	12	Y	Y	4
<code>int (*A4[3])</code>	Y	N	24	Y	N	8	Y	Y	4



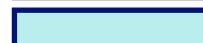
Allocated pointer



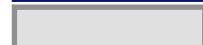
Unallocated pointer



Allocated int



Unallocated int



Understanding Pointers & Arrays #3

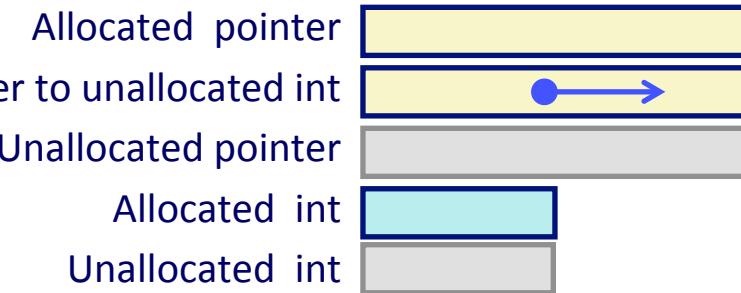
Decl	An			*An			**An		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
int A1[3][5]									
int *A2[3][5]									
int (*A3)[3][5]									
int *(A4[3][5])									
int (*A5[3])[5]									

Cmp: Compiles (Y/N)

**Bad: Possible bad
pointer reference (Y/
N)**

**Size: Value returned by
sizeof**

Decl	***An		
	Cmp	Bad	Size
int A1[3][5]			
int *A2[3][5]			
int (*A3)[3][5]			
int *(A4[3][5])			
int (*A5[3])[5]			



Declaration

`int A1[3][5]`

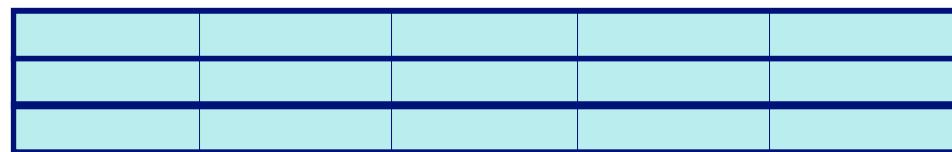
`int *A2[3][5]`

`int (*A3)[3][5]`

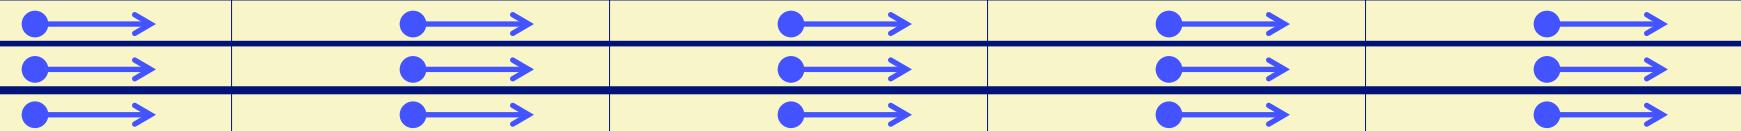
`int *(A4[3][5])`

`int (*A5[3])[5]`

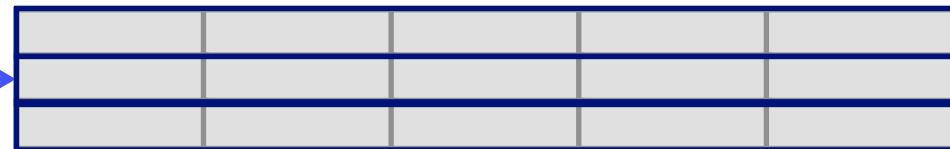
A1



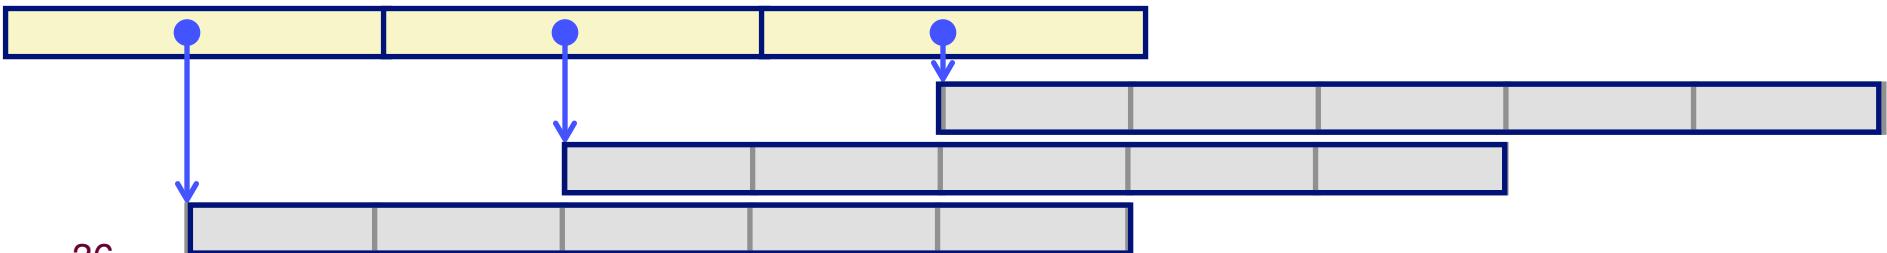
A2/A4



A3



A5



Understanding Pointers & Arrays #3

Decl	An			*An			**An		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
int A1[3][5]	Y	N	60	Y	N	20	Y	N	4
int *A2[3][5]	Y	N	120	Y	N	40	Y	N	8
int (*A3)[3][5]	Y	N	8	Y	Y	60	Y	Y	20
int *(A4[3][5])	Y	N	120	Y	N	40	Y	N	8
int (*A5[3])[5]	Y	N	24	Y	N	8	Y	Y	20

Cmp: Compiles (Y/N)

Bad: Possible bad pointer reference (Y/N)

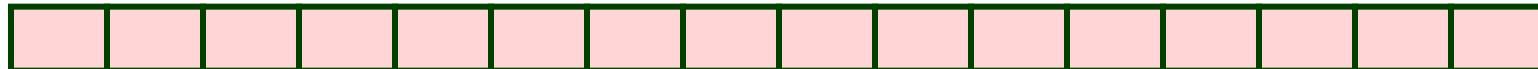
Size: Value returned by sizeof

Decl	***An		
	Cmp	Bad	Size
int A1[3][5]	N	-	-
int *A2[3][5]	Y	Y	4
int (*A3)[3][5]	Y	Y	4
int *(A4[3][5])	Y	Y	4
int (*A5[3])[5]	Y	Y	4

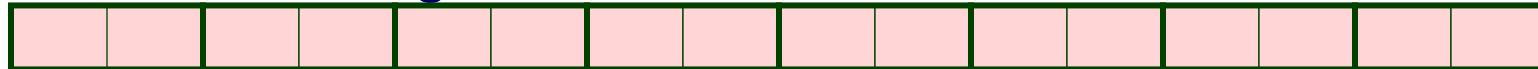
Aside: Programming with SSE3

XMM Registers

- 16 total, each 16 bytes
- 16 single-byte integers



- 8 16-bit integers



- 4 32-bit integers



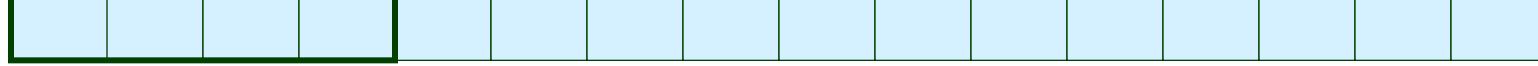
- 4 single-precision floats



- 2 double-precision floats



- 1 single-precision float



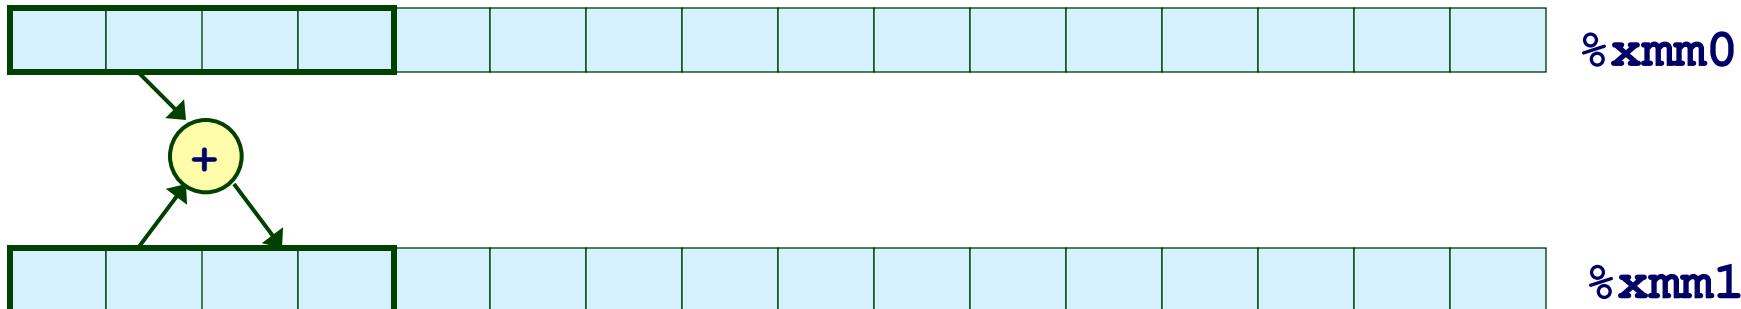
- 1 double-precision float



Aside: Scalar & SIMD Operations

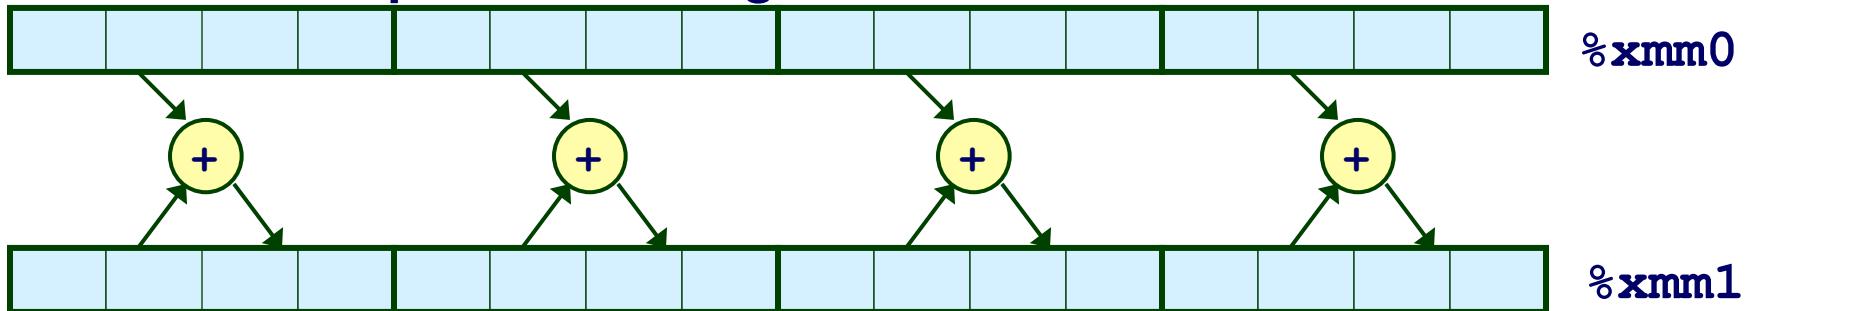
■ Scalar Operations: Single Precision

addss %xmm0 , %xmm1



■ SIMD Operations: Single Precision

addps %xmm0 , %xmm1



■ Scalar Operations: Double Precision

addsd %xmm0 , %xmm1

